Benjamin Chou

765-409-1859 | ben2002chou@gmail.com | linkedin.com/in/benjamin-chou-6aa058228 | github.com/ben2002chou | US Citizen

PROFESSIONAL SUMMARY

3rd-year ECE PhD Student at Purdue University specializing in AI/ML research for audio, music, and multimodal systems. Researching generative model steering, stem separation, Transformers for error detection, and training-free strategies for Multimodal LLMs. I am proficient in Audio DSP, Python, PyTorch, TensorFlow, and HPC clusters. I also have organizational and communication skills (experience organizing 3 AI4Music workshops).

SKILLS

- **LLM & Agent Systems:** LangChain, Vertex AI, prompt engineering, RAG, agent orchestration
- **Cloud & Developer Tooling:** Google Cloud Platform (GCP), GitHub CLI, GitHub API, Git, REST APIs
- **ML & Audio:** Python, PyTorch, TensorFlow, Transformers, Audio DSP, Linux, HPC

EDUCATION

Ph.D. in Electrical & Computer Engineering | *Purdue University*, West Lafayette, IN — 2023 - Present
- GPA: 4.0/4.0

B.Sc. in Electrical Engineering | *National Cheng Kung University*, Tainan, Taiwan — 2020 - 2023
- GPA: 4.0/4.3

EXPERIENCE

PhD Resident | *Google X*, Mountain View, CA — Apr '26 – Present
- Orchestrating LLM agents for code review workflows with **LangChain**, **Vertex AI**, **Google Cloud Platform (GCP)**, and **GitHub CLI**.

Audio Applied Research Scientist Intern (Incoming) | *Shure*, Niles, IL — Jul '26 – Sep '26

Applied Research Scientist Intern | *Reality Defender*, Manhattan, New York — May '25 – Sep '25
- Developed an in-context learning framework with large audio language models, achieving SOTA deepfake detection on in-the-wild speech (+54.9% avg Macro F1 improvement) (Natural Language Processing, RAG, ICL, Audio LLMs, and model routing).

Graduate Research Assistant | *Under Dr. Yung-Hsiang Lu, Purdue University*, West Lafayette, IN — Jan '24 – Present
- Developed an AI tutoring model with 6× higher F1 (AAAI 2025) and LadderSym with 2× better missed-note detection (ICLR 2026) (PyTorch, NLP, Audio Transformers).
- Leveraged HPC clusters and Linux environments to engineer generative pipelines for music mistakes, producing two datasets.
- Organized 'AI for Music' workshops at AAAI 2025, IEEE ICME 2025, and NeurIPS 2025, showcasing leadership and communication.

Startup Founder | *LocaLens*, West Lafayette, IN — Mar '23 – Jan '24
- Launched an app for menu translation, currency conversion, and allergen alerts using Flutter and Android Studio.
- Built an MVP in 4 months with weekly sprints, emphasizing leadership and collaboration.
- Integrated Google ML Toolkit for cross-platform AI with computer vision and real-time deployment.

Research Assistant | *Under Dr. Hsun-Ping Hsieh, National Cheng Kung University*, Taiwan — Mar '23 – Aug 2023
- Raised cancer-screening invitation accuracy to 83% using scikit-learn and tree-based classification (machine learning, statistics, data science).
- Elevated flash-chip yield by 17% with ensembling and Gaussian processes (NumPy, pandas).

PROJECTS

Improving Developer Code Understanding with GitHub Issues and Retrieval-Augmented Generation | *ChatGPT API, REST, Git*
- Built a RAG-based doc tool with GitHub issues; raised README clarity by 116% and code clarity by 16%.

Multi-Agent Self-Play for Beating Atari Games | *OpenAI Gym, PettingZoo, PyTorch, RL, Multi-agent systems*
- Used DQN/PPO self-play to improve Atari win rates by 11–32%.

Spectral Image Inpainting with Deep Learning | *MATLAB, C++, CUDA*
- Recovered 98% corrupted hyperspectral data at 0.02 RMSE using convex optimization and deep learning.

HONORS & AWARDS

- Outstanding Student Scholarship, awarded to top 2 students in NCKU-Purdue program (2020–2023).
- AAAI Student Scholarship (2025), selected for NeurIPS volunteer (2025).

PUBLICATIONS

- **B. Chou**, P. Jajal, N. J. Eliopoulos, J. C. Davis, G. K. Thiravathukal, K. Y.-J. Yun, Y.-H. Lu, "LadderSym: A Multimodal Interleaved Transformer for Music Practice Error Detection", *ICLR 2026 (top 3.85%)* [Link]
- **B. Chou**, et al., "Enabling Note-Wise Music Separation and Editing via a Partially Synthetic Data Engine", *In Progress*
- P. Jajal, N. J. Eliopoulos, **B. Chou**, G. K. Thiravathukal, J. C. Davis, Y. Lu, "AdaPerceiver: Transformers with Adaptive Width, Depth, and Tokens", *CVPR 2026 Findings* [Link]
- **B. Chou**, Y. Zhu, S. Koppisetti, "ICLAD: In-Context Learning with Comparison-Guidance for Audio Deepfake Detection", *ACL 2026 Findings*
- P. Jajal, N. J. Eliopoulos, **B. Chou**, G. K. Thiravathukal, J. C. Davis, Y. Lu, "Inference-Time Alignment of Diffusion Models with Evolutionary Algorithms", *CVPR 2026 Findings* [Link]
- **B. Chou**, P. Jajal, N. J. Eliopoulos, T. Nadolsky, C.-Y. Yang, N. Ravi, J. C. Davis, K. Y.-J. Yun, Y. Lu, "Detecting Performance Errors with Transformers", *AAAI 2025* [Link]
- P. Jajal, N. J. Eliopoulos, **B. Chou**, G. K. Thiravathukal, J. C. Davis, Y.-H. Lu, "Token Turing Machines are Efficient Vision Models", *WACV 2025* [Link]

PATENT PENDING

- System and Method for Detecting Musical Performance Errors, Patent Application 2025.
- ICLAD: In-Context Learning with Comparison-Guidance for Audio Deepfake Detection, Provisional Patent Application 2025.